

customer-sentiment-analysis

April 2, 2024

```
[1]: # all the necessary imports
import torch
from torch.utils.data import Dataset, DataLoader
import torch.nn as nn
from torch import optim
import torchtext
# from torchtext.data import Field, LabelField
# from torchtext.utils.data import TabularDataset
# from torchtext.utils.data import Iterator, BucketIterator
from torchtext.legacy.data import Field, TabularDataset, BucketIterator, \
    ↪Iterator
```

```
[2]: # set the seed
manual_seed = 572
torch.manual_seed(manual_seed)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

CNN's for text classification

Convolutional Neural Networks (CNNs) are essentially a special case of a normal feed forward network, instead of being “densely” connected (note you may see people refer to Feed Forward networks as “dense layers”), nodes in CNNs connect to a smaller set of nodes, defined by the “filters size” of the network. CNNs thus have a smaller local windows to look at data, but make up for it by generally using many additional filters, which might be able to learn different aspects of the data. These networks turn out to be extremely useful for processing images, audio, and anything with some sort of spatial properties to the data.

It turns out they can also be used for any sentence classification task. Words in a sentence it turns out have a sort of 1D spatial ordering, which means some classification tasks can benefit from this CNNs ability to operate over the length of the sequence. In addition, because of the sparsity of the connections, you end up being able to make much smaller networks that retain a great deal of power.

Pytorch 1D CNNs and max-pooling Here's a quick example of how these networks function:

```
[3]: x = torch.rand((2,5,10))    # batch size 2 with length 10 and 5 dim embedding.

in_dim = 5
filters = 4
```

```

cnn1d = nn.Conv1d(in_dim,filters,kernel_size=3,padding=1)
max_pool = nn.MaxPool1d(kernel_size=3, padding=1)
activation = nn.ReLU()
x = cnn1d(x)
x = activation(x)
print(x)
print("Take the highest value in each window using max pool")
print(max_pool(x))

```

```

tensor([[[[0.0115, 0.1652, 0.0000, 0.0210, 0.1613, 0.3903, 0.0000, 0.1766,
          0.1741, 0.2820],
          [0.0000, 0.0000, 0.0083, 0.1430, 0.0000, 0.0000, 0.2834, 0.0585,
          0.2041, 0.0000],
          [0.0456, 0.4565, 0.1949, 0.3690, 0.0888, 0.2262, 0.0732, 0.5068,
          0.5604, 0.3049],
          [0.0000, 0.0000, 0.0170, 0.1959, 0.0000, 0.0990, 0.0000, 0.0000,
          0.0000, 0.0000]],
         [[0.2387, 0.2856, 0.0000, 0.0000, 0.2042, 0.3228, 0.0630, 0.0000,
          0.0000, 0.2929],
          [0.0556, 0.2714, 0.0000, 0.0000, 0.0000, 0.1017, 0.1376, 0.1463,
          0.0000, 0.0000],
          [0.1951, 0.6214, 0.0341, 0.6379, 0.3405, 0.5166, 0.4090, 0.2046,
          0.3178, 0.5490],
          [0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
          0.0000, 0.2466]]], grad_fn=<ReluBackward0>)
Take the highest value in each window using max pool
tensor([[[[0.1652, 0.1613, 0.3903, 0.2820],
          [0.0000, 0.1430, 0.2834, 0.2041],
          [0.4565, 0.3690, 0.5068, 0.5604],
          [0.0000, 0.1959, 0.0990, 0.0000]],
         [[0.2856, 0.2042, 0.3228, 0.2929],
          [0.2714, 0.0000, 0.1463, 0.0000],
          [0.6214, 0.6379, 0.5166, 0.5490],
          [0.0000, 0.0000, 0.0000, 0.2466]]], grad_fn=<SqueezeBackward1>)

```

Here we've used a 1D CNN and max pooling to “summarize” our data, boiling a length 10 series to only 4 items.

```
[4]: cnn1d.weight.size()
```

```
[4]: torch.Size([4, 5, 3])
```

```
[5]: param_count = 0
for layer in cnn1d.parameters():
    param_count += layer.numel()
```

```
print(param_count)
```

64

At the first layer there's an embedding of length 5 which is passed through a CNN layer with kernel size of 3 and there are 4 such filters. Thus we get $5 \times 4 \times 3 = 60$ parameters. Finally there are 4 bias values, one for each filter. Hence the total number of parameters are 64.

Size through CNNs CNNs can be a little tricky because as the data passes through them the dimensions might change based on number of filters, padding, and two other things ,stride and dilation. MaxPooling also quickly can decrease the size of the data. This is useful especially as a way to “feature extract” or “dimensionality reduction” but it's important to make sure we get the right output dimensions.

```
[6]: ## Test your answer by making a random tensor of the appropriate size,  
## passing it through your proposed CNN+Maxpool and printing out the final size.
```

```
x = torch.rand((10,25,30))    # batch size 10 with length 30 and 25 dim  
↪ embedding.
```

```
in_dim = 25  
filters = 5
```

```
ks = 3
```

```
cnn1d = nn.Conv1d(in_dim,filters,kernel_size=ks,padding=ks//2)  
max_pool = nn.MaxPool1d(kernel_size=3, padding=1)  
activation = nn.ReLU()  
x = cnn1d(x)  
x = activation(x)
```

```
[7]: x.shape
```

```
[7]: torch.Size([10, 5, 30])
```

```
[8]: y = max_pool(x)  
y.shape
```

```
[8]: torch.Size([10, 5, 10])
```

1D CNN Model for Sentiment Analysis

```
[9]: import torchtext
```

```
# define the white space tokenizer to get tokens
```

```
def tokenize_en(tweet):
```

```
    """
```

```
    Tokenizes English tweet from a string into a list of strings (tokens)
```

```

        """
        return tweet.strip().split()

# define the TorchText's fields
TEXT = Field(sequential=True, tokenize=tokenize_en, lower=True)
LABEL = Field(sequential=False, unk_token = None)

train, val, test = TabularDataset.splits(
    path="./data/sentiment-twitter-2016-task4/", # the root directory where the
    ↪data lies
    train='train.tsv', validation="dev.tsv", test="test.tsv", # file names
    format='tsv',
    skip_header=False, # if your tsv file has a header, make sure to pass this
    ↪to ensure it doesn't get processed as data!
    fields=[('tweet', TEXT), ('label', LABEL)])

TEXT.build_vocab(train, min_freq=3) # builds vocabulary based on all the words
    ↪that occur at least twice in the training set
LABEL.build_vocab(train)

train_iter, val_iter, test_iter = BucketIterator.splits(
    (train, val, test), # we pass in the datasets we want the iterator to draw
    ↪data from
    batch_sizes=(64,64,64),
    sort_key=lambda x: len(x.tweet),
    sort=True,
    # A key to use for sorting examples in order to batch together examples with
    ↪similar lengths and minimize padding.
    sort_within_batch=True
)

VOCAB_SIZE = len(TEXT.vocab.stoi)
LABEL_SIZE = len(LABEL.vocab.stoi)

WORD_VEC_SIZE=300
# Note, the parameters to Embedding class below are:
# num_embeddings (int): size of the dictionary of embeddings
# embedding_dim (int): the size of each embedding vector
# For more details on Embedding class, see: https://github.com/pytorch/pytorch/
    ↪blob/master/torch/nn/modules/sparse.py

class ConvNet(nn.Module):

    def __init__(self, layer_num, filtersize, filters,nonlin, output_size,
    ↪VOCAB_SIZE, WORD_VEC_SIZE): #feel free to add additional parameters

```

```

super(ConvNet, self).__init__()
self.embedding = nn.Embedding(VOCAB_SIZE, WORD_VEC_SIZE, sparse=True)
self.embedding.weight.data.normal_(0.0,0.05)
self.layers = nn.ModuleList()
self.nonlin = nonlin
for i in range(layer_num):
    if i == 0:
        # YOUR CODE HERE    (FIRST LAYER CNN CODE)
        self.layers.append(nn.Conv1d(WORD_VEC_SIZE, filters,
↪kernel_size=filtersize, padding=filtersize//2))
    else:
        # YOUR CODE HERE    (LATER LAYER CNN CODE)
        self.layers.append(nn.Conv1d(filters,
↪filters,kernel_size=filtersize,padding=filtersize//2))
        self.layers.append(self.nonlin)

self.max_layer = nn.AdaptiveMaxPool1d(1)
self.output = nn.Linear(filters, output_size)
self.softmax = nn.LogSoftmax(dim=1)

def forward(self, x):
    # YOUR CODE HERE
    # PASS x THROUGH EMBEDDING (CHECK DIMENSIONS!!!):
    x = self.embedding(x) # Output: (L, N, C)
    # ENSURE DIMs CORRECT FOR CNN:
    x = x.permute(1, 2, 0) # CNN needs input (N,C,L)
    for layer in self.layers:
        x = layer(x)
    x = self.max_layer(x).squeeze(dim=-1)
    x =self.softmax(self.output(x))
    return x

```

```
[10]: from sklearn.metrics import accuracy_score
```

```

def train(loader,model,criterion,optimizer,device):
    total_loss = 0.0
    # iterate through the data loader
    num_sample = 0
    for batch in loader:
        # load the current batch
        batch_input = batch.tweet
        batch_output = batch.label

        batch_input = batch_input.to(device)
        batch_output = batch_output.to(device)
        # forward propagation
        # pass the data through the model

```

```

model_outputs = model(batch_input)
# compute the loss
cur_loss = criterion(model_outputs, batch_output)
total_loss += cur_loss.item()

# backward propagation (compute the gradients and update the model)
# clear the buffer
optimizer.zero_grad()
# compute the gradients
cur_loss.backward()
# update the weights
optimizer.step()

num_sample += batch_output.shape[0]
return total_loss/num_sample

# evaluation logic based on classification accuracy
def evaluate(loader,model,criterion,device):
    all_pred=[]
    all_label = []
    with torch.no_grad(): # impacts the autograd engine and deactivate it.
    ↪ reduces memory usage and speeds up computation
        for batch in loader:
            # load the current batch
            batch_input = batch.tweet
            batch_output = batch.label

            batch_input = batch_input.to(device)
            # forward propagation
            # pass the data through the model
            model_outputs = model(batch_input)
            # identify the predicted class for each example in the batch
            probabilities, predicted = torch.max(model_outputs.cpu().data, 1)
            # put all the true labels and predictions to two lists
            all_pred.extend(predicted)
            all_label.extend(batch_output)

    accuracy = accuracy_score(all_label, all_pred)
    return accuracy

```

Example run through:

```

[11]: model = ConvNet(4, 3, 10, nn.ReLU(), LABEL_SIZE, VOCAB_SIZE, WORD_VEC_SIZE)
      model = model.to(device)
      model

```

```
[11]: ConvNet(
  (embedding): Embedding(3330, 300, sparse=True)
  (layers): ModuleList(
    (0): Conv1d(300, 10, kernel_size=(3,), stride=(1,), padding=(1,))
    (1): ReLU()
    (2): Conv1d(10, 10, kernel_size=(3,), stride=(1,), padding=(1,))
    (3): ReLU()
    (4): Conv1d(10, 10, kernel_size=(3,), stride=(1,), padding=(1,))
    (5): ReLU()
    (6): Conv1d(10, 10, kernel_size=(3,), stride=(1,), padding=(1,))
    (7): ReLU()
  )
  (nonlin): ReLU()
  (max_layer): AdaptiveMaxPool1d(output_size=1)
  (output): Linear(in_features=10, out_features=3, bias=True)
  (softmax): LogSoftmax(dim=1)
)
```

```
[12]: # Getting a batch of data

for batch in train_iter:
    tweets = batch.tweet
    labels = batch.label
    break #we use first batch as an example.
```

```
[13]: tweets.shape
```

```
[13]: torch.Size([8, 64])
```

```
[14]: # Passing input to GPU
tweets = tweets.to(device)
labels = labels.to(device)

out = model(tweets)
out
```

```
[14]: tensor([[ -1.3135, -0.8572, -1.1816],
          [ -1.3131, -0.8567, -1.1827],
          [ -1.3132, -0.8567, -1.1826],
          [ -1.3135, -0.8567, -1.1823],
          [ -1.3136, -0.8566, -1.1825],
          [ -1.3135, -0.8565, -1.1827],
          [ -1.3136, -0.8569, -1.1820],
          [ -1.3133, -0.8575, -1.1814],
          [ -1.3135, -0.8572, -1.1816],
          [ -1.3139, -0.8575, -1.1809],
          [ -1.3134, -0.8571, -1.1818],
```

[-1.3137, -0.8569, -1.1819],
 [-1.3133, -0.8571, -1.1819],
 [-1.3137, -0.8576, -1.1809],
 [-1.3135, -0.8568, -1.1822],
 [-1.3137, -0.8573, -1.1814],
 [-1.3135, -0.8571, -1.1817],
 [-1.3137, -0.8573, -1.1814],
 [-1.3137, -0.8572, -1.1814],
 [-1.3135, -0.8566, -1.1824],
 [-1.3129, -0.8571, -1.1823],
 [-1.3136, -0.8573, -1.1815],
 [-1.3137, -0.8573, -1.1813],
 [-1.3135, -0.8569, -1.1820],
 [-1.3134, -0.8572, -1.1817],
 [-1.3136, -0.8573, -1.1815],
 [-1.3136, -0.8567, -1.1822],
 [-1.3137, -0.8566, -1.1823],
 [-1.3135, -0.8570, -1.1820],
 [-1.3136, -0.8568, -1.1821],
 [-1.3135, -0.8568, -1.1822],
 [-1.3144, -0.8575, -1.1804],
 [-1.3131, -0.8572, -1.1821],
 [-1.3137, -0.8570, -1.1817],
 [-1.3138, -0.8575, -1.1810],
 [-1.3140, -0.8566, -1.1821],
 [-1.3133, -0.8574, -1.1815],
 [-1.3129, -0.8573, -1.1820],
 [-1.3134, -0.8564, -1.1829],
 [-1.3139, -0.8570, -1.1816],
 [-1.3134, -0.8565, -1.1827],
 [-1.3136, -0.8570, -1.1819],
 [-1.3134, -0.8573, -1.1816],
 [-1.3132, -0.8576, -1.1813],
 [-1.3137, -0.8568, -1.1821],
 [-1.3137, -0.8575, -1.1810],
 [-1.3140, -0.8566, -1.1821],
 [-1.3136, -0.8573, -1.1813],
 [-1.3134, -0.8576, -1.1812],
 [-1.3136, -0.8572, -1.1816],
 [-1.3135, -0.8572, -1.1816],
 [-1.3135, -0.8572, -1.1816],
 [-1.3139, -0.8570, -1.1816],
 [-1.3132, -0.8579, -1.1809],
 [-1.3132, -0.8577, -1.1812],
 [-1.3132, -0.8581, -1.1806],
 [-1.3134, -0.8578, -1.1810],
 [-1.3135, -0.8576, -1.1811],


```

[-1.3134, -0.8568, -1.1822],
[-1.3134, -0.8568, -1.1822],
[-1.3133, -0.8575, -1.1813],
[-1.3137, -0.8572, -1.1815],
[-1.3133, -0.8577, -1.1811],
[-1.3134, -0.8575, -1.1814]], grad_fn=<LogSoftmaxBackward0>)

```

```

[15]: LEARNING_RATE = 0.01
      criterion = nn.NLLLoss()
      optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)

```

```

[16]: loss = criterion(out, labels)
      loss

```

```

[16]: tensor(1.1290, grad_fn=<NllLossBackward0>)

```

```

[17]: # Sanity code to check if all works
      train(train_iter, model, criterion, optimizer, device)

```

```

[17]: 0.01713936573266983

```

```

[18]: evaluate(train_iter, model, criterion, device)

```

```

[18]: 0.3405

```

```

[19]: evaluate(val_iter, model, criterion, device)

```

```

[19]: 0.3826913456728364

```

```

[20]: evaluate(test_iter, model, criterion, device)

```

```

[20]: 0.5012601783637068

```

1D CNN Performance Based on our initial network we'd like to compare how depth matters vs number of filters in a given layer.

```

[21]: import scipy.stats

      LEARNING_RATE=.1
      MAX_EPOCHS=10

      def random_search(num_iter):
          results = []
          for i in range(num_iter):
              config = {
                  #define hyperparameters here
                  "layers": scipy.stats.randint.rvs(1,3),

```

```

        "filters": scipy.stats.randint.rvs(10,200)
    }

    print("new config")
    print(config)
    model = ConvNet(config["layers"],3,config["filters"],nn.
↪ReLU(),output_size=3, VOCAB_SIZE=VOCAB_SIZE, WORD_VEC_SIZE=WORD_VEC_SIZE)
    model.to(device)
    criterion = nn.NLLLoss()
    optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)

    max_val = 0
    best_epoch = 0
    for epoch in range(MAX_EPOCHS):
        # train the model for one pass over the data
        train_loss = train(train_iter,model,criterion,optimizer,device)
        # compute the training accuracy
        train_acc = evaluate(train_iter,model,criterion,device)
        # compute the validation accuracy
        val_acc = evaluate(val_iter,model,criterion,device)
        if val_acc > max_val:
            max_val = val_acc
            best_epoch = epoch+1
        # print the loss for every epoch
        print('Epoch [{}/{}], Loss: {:.4f}, Training Accuracy: {:.4f},
↪Validation Accuracy: {:.4f}'.format(epoch+1, MAX_EPOCHS, train_loss,
↪train_acc, val_acc))
        results.append((max_val,best_epoch,config))
    return results

```

[22]: random_search(20)

```

new config
{'layers': 1, 'filters': 148}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5155, Validation Accuracy:
0.4222
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5168, Validation Accuracy:
0.4302

```

Epoch [7/10], Loss: 0.0150, Training Accuracy: 0.5258, Validation Accuracy: 0.4452
Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5445, Validation Accuracy: 0.4622
Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5615, Validation Accuracy: 0.4637
Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5727, Validation Accuracy: 0.4642
new config
{'layers': 1, 'filters': 184}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5155, Validation Accuracy: 0.4212
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5167, Validation Accuracy: 0.4292
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5263, Validation Accuracy: 0.4442
Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5430, Validation Accuracy: 0.4622
Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5600, Validation Accuracy: 0.4682
Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5747, Validation Accuracy: 0.4652
new config
{'layers': 1, 'filters': 144}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4222
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5180, Validation Accuracy: 0.4277
Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5285, Validation Accuracy: 0.4437

Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5522, Validation Accuracy: 0.4587
Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5743, Validation Accuracy: 0.4667
new config
{'layers': 1, 'filters': 89}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5157, Validation Accuracy: 0.4232
Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5220, Validation Accuracy: 0.4362
Epoch [9/10], Loss: 0.0149, Training Accuracy: 0.5403, Validation Accuracy: 0.4477
Epoch [10/10], Loss: 0.0147, Training Accuracy: 0.5598, Validation Accuracy: 0.4632
new config
{'layers': 1, 'filters': 75}
Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5178, Validation Accuracy: 0.4287
Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5247, Validation Accuracy: 0.4432
Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5430, Validation Accuracy: 0.4567
Epoch [10/10], Loss: 0.0147, Training Accuracy: 0.5625, Validation Accuracy: 0.4677

```

new config
{'layers': 2, 'filters': 173}
Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [2/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [5/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [6/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [7/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [8/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [9/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [10/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
new config
{'layers': 1, 'filters': 73}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [6/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [7/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy:
0.4222
Epoch [8/10], Loss: 0.0151, Training Accuracy: 0.5183, Validation Accuracy:
0.4302
Epoch [9/10], Loss: 0.0150, Training Accuracy: 0.5375, Validation Accuracy:
0.4432
Epoch [10/10], Loss: 0.0148, Training Accuracy: 0.5600, Validation Accuracy:
0.4502
new config
{'layers': 1, 'filters': 188}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy:
0.4217

```

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5155, Validation Accuracy: 0.4237

Epoch [7/10], Loss: 0.0150, Training Accuracy: 0.5218, Validation Accuracy: 0.4387

Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5363, Validation Accuracy: 0.4557

Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5593, Validation Accuracy: 0.4682

Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5780, Validation Accuracy: 0.4727

new config

{'layers': 1, 'filters': 197}

Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5155, Validation Accuracy: 0.4227

Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5180, Validation Accuracy: 0.4317

Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5310, Validation Accuracy: 0.4467

Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5537, Validation Accuracy: 0.4652

Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5750, Validation Accuracy: 0.4697

new config

{'layers': 2, 'filters': 175}

Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [7/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [8/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [9/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [10/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

new config

{'layers': 1, 'filters': 165}

Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5155, Validation Accuracy: 0.4222

Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5195, Validation Accuracy: 0.4297

Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5315, Validation Accuracy: 0.4557

Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5548, Validation Accuracy: 0.4717

Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5757, Validation Accuracy: 0.4822

new config

{'layers': 1, 'filters': 110}

Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4222

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5167, Validation Accuracy: 0.4282

Epoch [7/10], Loss: 0.0150, Training Accuracy: 0.5255, Validation Accuracy: 0.4492

Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5453, Validation Accuracy: 0.4562

Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5615, Validation Accuracy: 0.4512

Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5788, Validation Accuracy: 0.4647

new config

```
{'layers': 2, 'filters': 97}
```

Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [7/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [8/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [9/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [10/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

new config

```
{'layers': 1, 'filters': 99}
```

Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5160, Validation Accuracy: 0.4292

Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5238, Validation Accuracy: 0.4367

Epoch [9/10], Loss: 0.0149, Training Accuracy: 0.5380, Validation Accuracy: 0.4597

Epoch [10/10], Loss: 0.0147, Training Accuracy: 0.5560, Validation Accuracy: 0.4617

new config

{'layers': 1, 'filters': 104}

Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5155, Validation Accuracy: 0.4217

Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5180, Validation Accuracy: 0.4247

Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5280, Validation Accuracy: 0.4397

Epoch [9/10], Loss: 0.0149, Training Accuracy: 0.5475, Validation Accuracy: 0.4597

Epoch [10/10], Loss: 0.0147, Training Accuracy: 0.5647, Validation Accuracy: 0.4682

new config

{'layers': 1, 'filters': 24}

Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4227

Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5190, Validation Accuracy: 0.4257

Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5273, Validation Accuracy: 0.4412

Epoch [9/10], Loss: 0.0149, Training Accuracy: 0.5427, Validation Accuracy: 0.4482

```

Epoch [10/10], Loss: 0.0148, Training Accuracy: 0.5547, Validation Accuracy:
0.4557
new config
{'layers': 1, 'filters': 129}
Epoch [1/10], Loss: 0.0157, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [4/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5155, Validation Accuracy:
0.4222
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5178, Validation Accuracy:
0.4282
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5260, Validation Accuracy:
0.4422
Epoch [8/10], Loss: 0.0149, Training Accuracy: 0.5412, Validation Accuracy:
0.4602
Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5570, Validation Accuracy:
0.4622
Epoch [10/10], Loss: 0.0146, Training Accuracy: 0.5720, Validation Accuracy:
0.4607
new config
{'layers': 2, 'filters': 105}
Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [2/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [3/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [5/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [6/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [7/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [8/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [9/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
Epoch [10/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy:
0.4217
new config
{'layers': 1, 'filters': 59}

```

Epoch [1/10], Loss: 0.0159, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0153, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [6/10], Loss: 0.0152, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [7/10], Loss: 0.0151, Training Accuracy: 0.5168, Validation Accuracy: 0.4262
Epoch [8/10], Loss: 0.0150, Training Accuracy: 0.5295, Validation Accuracy: 0.4502
Epoch [9/10], Loss: 0.0148, Training Accuracy: 0.5568, Validation Accuracy: 0.4582
Epoch [10/10], Loss: 0.0147, Training Accuracy: 0.5732, Validation Accuracy: 0.4707

new config

```
{'layers': 2, 'filters': 147}
```

Epoch [1/10], Loss: 0.0158, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [2/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [3/10], Loss: 0.0155, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [4/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [5/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [6/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [7/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [8/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [9/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217
Epoch [10/10], Loss: 0.0154, Training Accuracy: 0.5157, Validation Accuracy: 0.4217

[22]: [(0.464232116058029, 10, {'layers': 1, 'filters': 148}),
(0.46823411705852924, 9, {'layers': 1, 'filters': 184}),
(0.4667333668334164, 10, {'layers': 1, 'filters': 144}),
(0.46323161580790395, 10, {'layers': 1, 'filters': 89}),

```
(0.46773386693346675, 10, {'layers': 1, 'filters': 75}),
(0.42171085542771386, 1, {'layers': 2, 'filters': 173}),
(0.4502251125562781, 10, {'layers': 1, 'filters': 73}),
(0.47273636818409204, 10, {'layers': 1, 'filters': 188}),
(0.46973486743371684, 10, {'layers': 1, 'filters': 197}),
(0.42171085542771386, 1, {'layers': 2, 'filters': 175}),
(0.4822411205602801, 10, {'layers': 1, 'filters': 165}),
(0.46473236618309155, 10, {'layers': 1, 'filters': 110}),
(0.42171085542771386, 1, {'layers': 2, 'filters': 97}),
(0.46173086543271635, 10, {'layers': 1, 'filters': 99}),
(0.46823411705852924, 10, {'layers': 1, 'filters': 104}),
(0.45572786393196596, 10, {'layers': 1, 'filters': 24}),
(0.4622311155577789, 9, {'layers': 1, 'filters': 129}),
(0.42171085542771386, 1, {'layers': 2, 'filters': 105}),
(0.47073536768384194, 10, {'layers': 1, 'filters': 59}),
(0.42171085542771386, 1, {'layers': 2, 'filters': 147})]
```

1D CNN summary My best score is 48.47% with 2 layers and 73 filters. Based on the results having two layers gave a better score and having filters in the range of 50-80 gave the best results, anything higher reduced the scores possibly due to overfitting.

0.1 Building a Recurrent Neural Network

we use a corpus from the [CL-Aff shared task](#). HappyDB is a dataset of about 100,000 happy moments crowd-sourced via Amazon’s Mechanical Turk where each worker was asked to describe in a complete sentence **what made them happy in the past 24 hours**. Each user was asked to describe three such moments.

We have already preprocessed (tokenization, removing URLs, mentions, hashtags and so on) the tweets and placed it under `data/happy_db` folder in three files as `train.tsv`, `dev.tsv` and `test.tsv`.

Whitespace tokenizer

```
[40]: from nltk.tokenize import WhitespaceTokenizer

def whitespace_tokenize(text):

    # your code goes here
    tokens = WhitespaceTokenizer().tokenize(text)
    return tokens
```

TorchText’s Fields

```
[41]: # your code goes here
TEXT = Field(sequential=True, tokenize=whitespace_tokenize, lower=False)
LABEL = Field(sequential=False, unk_token = None)
```

TabularDataset class and Fields

```
[42]: # your code goes here
train, val, test = TabularDataset.splits(path= "./data/happy_db/", train='train.
↳tsv', validation="dev.tsv", test="test.tsv", # file names
    format='tsv',
    skip_header=True, # if your tsv file has a header, make sure to pass this
↳to ensure it doesn't get processed as data!
    fields=[('tweet', TEXT), ('label', LABEL)])
```

Building vocab

```
[43]: # your code goes here
TEXT.build_vocab(train, max_size=5000) # builds vocabulary based on all the
↳words that occur at least twice in the training set
LABEL.build_vocab(train, max_size = 5000)
```

```
[44]: # your code goes here
print(len(TEXT.vocab.stoi))
print(len(LABEL.vocab.stoi))
```

5002

2

Constructing Iterators

```
[45]: # your code goes here
# from torchtext.data import Iterator, BucketIterator

train_iter, val_iter, test_iter = BucketIterator.splits(
    (train, val, test), # we pass in the datasets we want the iterator to draw
↳data from
    batch_sizes=(32,32,32),
    sort_key=lambda x: len(x.tweet),
    sort=True,
    # A key to use for sorting examples in order to batch together examples with
↳similar lengths and minimize padding.
    sort_within_batch=True
)
```

Model creation

```
[46]: import torch.nn as nn
class LSTMmodel(nn.Module):

    def __init__(self, embedding_size, vocab_size, output_size, hidden_size,
↳num_layers):
    # In the constructor we define the layers for our model
    super(LSTMmodel, self).__init__()
```

```

    # word embedding lookup table
    self.embedding = nn.Embedding(num_embeddings=vocab_size,
    ↪embedding_dim=embedding_size, sparse=True)

    # core LSTM module
    self.LSTM_layer = nn.LSTM(input_size=300, hidden_size=500, num_layers=2)
    # activation function
    self.activation_fn = nn.Tanh()
    # classification related modules
    self.linear_layer = nn.Linear(hidden_size, output_size)
    self.softmax_layer = nn.LogSoftmax(dim=1)
    self.debug = False

def forward(self, x):
    # In the forward function we define the forward propagation logic
    if self.debug:
        print("input word indices shape = ", x.size())
    out = self.embedding(x)
    if self.debug:
        print("word embeddings shape = ", out.size())
    out, _ = self.LSTM_layer(out) # since we are not feeding h_0 explicitly,
    ↪h_0 will be initialized to zeros by default
    if self.debug:
        print("RNN output (features from last layer of RNN for all timesteps)
    ↪shape = ", out.size())
    # classify based on the hidden representation after RNN processes the last
    ↪token
    out = out[-1]
    if self.debug:
        print("Tweet embeddings or RNN output (features from last layer of RNN
    ↪for the last timestep only) shape = ", out.size())
    out = self.activation_fn(out)
    if self.debug:
        print("ReLU output shape = ", out.size())
    out = self.linear_layer(out)
    if self.debug:
        print("linear layer output shape = ", out.size())
    out = self.softmax_layer(out) # accepts 2D or more dimensional inputs
    if self.debug:
        print("softmax layer output shape = ", out.size())
    return out

```

Instantiating the model

```

[47]: # your code goes here
HIDDEN_SIZE = 500 # no. of units in the hidden layer
NUM_LAYERS = 2
EMBEDDING_SIZE = 300

```

```
VOCAB_SIZE = len(TEXT.vocab.stoi)
NUM_CLASSES = len(LABEL.vocab.stoi)
```

```
[48]: model = LSTMmodel(EMBEDDING_SIZE, VOCAB_SIZE, NUM_CLASSES, HIDDEN_SIZE,
    ↪NUM_LAYERS)
print(model)
#model.to(device)
```

```
LSTMmodel(
  (embedding): Embedding(5002, 300, sparse=True)
  (LSTM_layer): LSTM(300, 500, num_layers=2)
  (activation_fn): Tanh()
  (linear_layer): Linear(in_features=500, out_features=2, bias=True)
  (softmax_layer): LogSoftmax(dim=1)
)
```

Create an optimizer for training

```
[49]: LEARNING_RATE = 0.1
criterion = nn.NLLLoss()
# create an instance of SGD with required hyperparameters
optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)
```

```
[50]: # your code goes here
# your code goes here
total_parameters = 0
for variable in model.parameters():
    # shape is an array of tf.Dimension
    shape = variable.shape
    variable_parameters = 1
    for dim in shape:
        variable_parameters *= dim
    total_parameters += variable_parameters
print("Total Parameters:",total_parameters)
```

Total Parameters: 5109602

```
[51]: points = 5110103
bits_per_point = 32
number_of_mega_byte= (((points*bits_per_point)/8)/10**6)
print("Model Mega Byte memory:",number_of_mega_byte)
```

Model Mega Byte memory: 20.440412

training and evaluation

```
[52]: import numpy as np
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
```

```

def train(loader):
    total_loss = 0.0
    # iterate through the data loader
    num_sample = 0
    for batch in loader:
        # load the current batch
        batch_input = batch.tweet
        batch_output = batch.label

        batch_input = batch_input.to(device)
        batch_output = batch_output.to(device)
        # forward propagation
        # pass the data through the model
        model_outputs = model(batch_input)
        # compute the loss
        cur_loss = criterion(model_outputs, batch_output)
        total_loss += cur_loss.item()

        # backward propagation (compute the gradients and update the model)
        # clear the buffer
        optimizer.zero_grad()
        # compute the gradients
        cur_loss.backward()
        # update the weights
        optimizer.step()

    num_sample += batch_output.shape[0]
    return total_loss/num_sample

# evaluation logic based on classification accuracy
def evaluate(loader):
    all_pred=[]
    all_label = []
    with torch.no_grad(): # impacts the autograd engine and deactivate it. ↵
    ↪ reduces memory usage and speeds up computation
    for batch in loader:
        # load the current batch
        batch_input = batch.tweet
        batch_output = batch.label

        batch_input = batch_input.to(device)
        # forward propagation
        # pass the data through the model
        model_outputs = model(batch_input)
        # identify the predicted class for each example in the batch
        probabilities, predicted = torch.max(model_outputs.cpu().data, 1)

```



```

        # put all the true labels and predictions to two lists
        all_pred.extend(predicted)
        all_label.extend(batch_output)

    accuracy = accuracy_score(all_label, all_pred)
    f1score = f1_score(all_label, all_pred, average='macro')
    return accuracy, f1score

```

Saving the model

```

[53]: # your code goes here
      # Print model's state_dict
      import os
      if not os.path.exists("./ckpt"): # check if the directory doesn't exist already
          os.mkdir("./ckpt")

```

```

[54]: # start the training
      MAX_EPOCHS = 10
      for epoch in range(MAX_EPOCHS):
          # train the model for one pass over the data
          train_loss = train(train_iter)
          # compute the training accuracy
          train_acc = evaluate(train_iter)
          # compute the validation accuracy
          val_acc = evaluate(val_iter)

          # print the loss for every epoch
          print('epoch ', epoch+1, 'loss ', train_loss, 'Train Accuracy & F1', train_acc, 'Validation Accuracy & F1 ', val_acc)

          # save model, optimizer, and number of epoch to a dictionary
          model_save = {
              'epoch': epoch, # number of epoch
              'model_state_dict': model.state_dict(), # model parameters
              'optimizer_state_dict': optimizer.state_dict(), # save optimizer
              'loss': train_loss # training loss
          }

          # use torch.save to store
          torch.save(model_save, "./ckpt/model_{}.pt".format(epoch))

```

```

epoch 1 loss 0.019151628931109426 Train Accuracy & F1 (0.7309422348484849,
0.7137388123133535) Validation Accuracy & F1 (0.6732954545454546,
0.6316692464243143)
epoch 2 loss 0.013246820703374617 Train Accuracy & F1 (0.8325047348484849,
0.8324948186057661) Validation Accuracy & F1 (0.7926136363636364,
0.7924571141394334)

```

```

epoch 3 loss 0.011463872385458231 Train Accuracy & F1 (0.7906013257575758,
0.7885851163107714) Validation Accuracy & F1 (0.7679924242424242,
0.7676421588658725)
epoch 4 loss 0.010660828917576564 Train Accuracy & F1 (0.8290719696969697,
0.8284218987839711) Validation Accuracy & F1 (0.7897727272727273,
0.7894396482038348)
epoch 5 loss 0.009733532781176495 Train Accuracy & F1 (0.8325047348484849,
0.8315165046642996) Validation Accuracy & F1 (0.7878787878787878,
0.7867151404983717)
epoch 6 loss 0.008951072622122329 Train Accuracy & F1 (0.8135653409090909,
0.8113326949609299) Validation Accuracy & F1 (0.7471590909090909,
0.7425695386005999)
epoch 7 loss 0.008422675708337038 Train Accuracy & F1 (0.8693181818181818,
0.8691640062979097) Validation Accuracy & F1 (0.8210227272727273,
0.8210225667744264)
epoch 8 loss 0.007800204066425618 Train Accuracy & F1 (0.8938210227272727,
0.8937999504810787) Validation Accuracy & F1 (0.8248106060606061,
0.8246783840903891)
epoch 9 loss 0.0071551693174776365 Train Accuracy & F1 (0.8894412878787878,
0.889384151593454) Validation Accuracy & F1 (0.8172348484848485,
0.816931303428496)
epoch 10 loss 0.0065436911794346415 Train Accuracy & F1 (0.8776041666666666,
0.8752106220183131) Validation Accuracy & F1 (0.8418560606060606,
0.8366226150275844)

```

```

[55]: # define a new model

model2 = LSTMmodel(EMBEDDING_SIZE, VOCAB_SIZE, NUM_CLASSES, HIDDEN_SIZE,
    ↪NUM_LAYERS)

# load checkpoint

checkpoint = torch.load("./ckpt/model_9.pt")

# assign the parameters of checkpoint to this new model

model2.load_state_dict(checkpoint['model_state_dict'])
model2.to(device)

print(model2) # can be used for inference or for further training

```

```

LSTMmodel(
  (embedding): Embedding(5002, 300, sparse=True)
  (LSTM_layer): LSTM(300, 500, num_layers=2)
  (activation_fn): Tanh()
  (linear_layer): Linear(in_features=500, out_features=2, bias=True)
  (softmax_layer): LogSoftmax(dim=1)
)

```

Model Evaluation

```
[56]: # your code goes here
      accuracy, f1score = evaluate(test_iter)
      print("Test Accuracy : ", accuracy)
      print("Test F1 : ", f1score)
```

Test Accuracy : 0.8314393939393939

Test F1 : 0.8265100807419248